

UW Datathon 2026 Preprocessing

Kirby Vo, Kelbin Luong

2026-02-07

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## vforcats   1.0.0     v stringr   1.5.1
## v ggplot2   3.5.1     v tibble    3.2.1
## v lubridate 1.9.4     v tidyrr    1.3.1
## v purrr    1.0.2

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(stringr)
library(writexl)
```

LOAD DATA

```
data <- read_delim("datasets/Access_to_Tech_Dataset.csv")

## Rows: 3524 Columns: 17
## -- Column specification -----
## Delimiter: ","
## chr (14): id, domain_category, web_URL, scrape_status, html_file_name, html_...
## dbl  (3): web_URL_id, violation_count, violation_score
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

PREVIEW DATA

```
data |>
  names()
```

```

## [1] "id"                               "web_URL_id"
## [3] "domain_category"                  "web_URL"
## [5] "scrape_status"                   "html_file_name"
## [7] "html_file_path"                  "violation_count"
## [9] "violation_name"                 "violation_score"
## [11] "violation_description"          "violation_description_url"
## [13] "affected_html_elements"         "violation_category"
## [15] "violation_impact"              "wcag_reference"
## [17] "supplementary_information"

data |>
  sample_n(10)

## # A tibble: 10 x 17
##   id      web_URL_id domain_category  web_URL scrape_status html_file_name
##   <chr>    <dbl> <chr>           <chr>   <chr>        <chr>
## 1 6699878654_4 6699878654 News and Media https://www.propublica.org
## 2 5117799549_8 5117799549 E-commerce     https://www.alibabacloud.com
## 3 193670859_12 193670859 News and Media  https://www.abcnews.go.com
## 4 1851235277_1 1851235277 E-commerce     https://www.rakuten.com
## 5 501_1          501 Health and Wellness https://www.healthcare.com
## 6 7124602025_2 7124602025 News and Media https://www.bostonglobe.com
## 7 8056970126_5 8056970126 Streaming Platform https://www.grokkr.com
## 8 1488388154_2 1488388154 Technology Science https://www.cnet.com
## 9 503_6          503 Technology Science https://arstechnica.com
## 10 8866935042_2 8866935042 Streaming Platform https://www.fox.com

## # i 11 more variables: html_file_path <chr>, violation_count <dbl>,
## #   violation_name <chr>, violation_score <dbl>, violation_description <chr>,
## #   violation_description_url <chr>, affected_html_elements <chr>,
## #   violation_category <chr>, violation_impact <chr>, wcag_reference <chr>,
## #   supplementary_information <chr>

```

SEARCH FOR DATA INCONSISTENCIES

We found that looking at distinct categories, we saw some inconsistent domain category entries that we would need to remap

```

data |>
  distinct(domain_category)

## # A tibble: 9 x 1
##   domain_category
##   <chr>
## 1 Government and Public Services
## 2 News and Media
## 3 Technology Science and Research
## 4 E-commerce
## 5 Educational Platforms
## 6 Streaming Platforms
## 7 Health and Wellness
## 8 TechnologyScienceResearch
## 9 Ecommerce

```

We also found cases where if a url had multiple domain categories assigned to it, any violation entries for it would be duplicated accordingly.

```
data |>
  filter(web_URL == "https://arstechnica.com/health/") |>
  select(domain_category, web_URL, violation_name, violation_score) |>
  head(2)
```

```
## # A tibble: 2 x 4
##   domain_category      web_URL      violation_name violation_score
##   <chr>          <chr>          <chr>                  <dbl>
## 1 News and Media https://arste~ color-contras~           4
## 2 Technology Science and Research https://arste~ color-contras~           4
```

We had also noticed while processing the data that there were instances of urls with or without /, which would count them as separate webpages despite being the same, which needed to be addressed.

```
data |>
  filter(web_URL == "https://www.nbcnews.com" | web_URL == "https://www.nbcnews.com/") |>
  select(web_URL) |>
  slice(1, 22)

## # A tibble: 2 x 1
##   web_URL
##   <chr>
## 1 https://www.nbcnews.com
## 2 https://www.nbcnews.com/
```

Regarding domains, we found that there were a lot of domains such as arstechnica.com which had multiple webpages under it, which meant we would need to group them properly by domain in order to answer domain specific questions.

```
data |>
  filter(str_detect(web_URL, "arstechnica\\.com")) |>
  distinct(web_URL) |>
  head(5)

## # A tibble: 5 x 1
##   web_URL
##   <chr>
## 1 https://arstechnica.com/ai/
## 2 https://arstechnica.com/information-technology/
## 3 https://arstechnica.com/cars/
## 4 https://arstechnica.com/culture/
## 5 https://arstechnica.com/gaming/
```

CLEAN DATA

Here we create a map to convert weird domain category values to correct ones.

```

renaming_map <- c(
  "Government and Public Services" = "Government and Public Services",
  "News and Media" = "News and Media",
  "Technology Science and Research" = "Technology Science and Research",
  "E-commerce" = "E-commerce",
  "Educational Platforms" = "Educational Platforms",
  "Streaming Platforms" = "Streaming Platforms",
  "Health and Wellness" = "Health and Wellness",
  "TechnologyScienceResearch" = "Technology Science and Research",
  "Ecommerce" = "E-commerce"
)

data_clean <- data |>
  mutate(
    domain_category = renaming_map[domain_category],
    domain = web_URL |>
      str_replace("https?://(www\\.)?", "") |>
      str_replace("/.*$", "")
  )

```

Ensure that violations aren't duplicated based off number of additional domain categories.

```

data_clean_deduped <- data_clean |>
  distinct(web_URL,
           violation_name,
           violation_score,
           violation_category,
           violation_impact,
           wcag_reference,
           .keep_all = TRUE)

```

VERIFY PROBLEMS FIXED

```

data_clean_deduped |>
  distinct(domain_category)

## # A tibble: 7 x 1
##   domain_category
##   <chr>
## 1 Government and Public Services
## 2 News and Media
## 3 E-commerce
## 4 Technology Science and Research
## 5 Educational Platforms
## 6 Streaming Platforms
## 7 Health and Wellness

data_clean_deduped |>
  filter(web_URL == "https://www.nbcnews.com/" | web_URL == "https://www.nbcnews.com") |>
  select(web_URL, domain) |>
  slice(c(1, 22))

```

```

## # A tibble: 2 x 2
##   web_URL           domain
##   <chr>             <chr>
## 1 https://www.nbcnews.com nbcnews.com
## 2 https://www.nbcnews.com/ nbcnews.com

data_clean_deduped |>
  filter(web_URL == "https://arstechnica.com/health/") |>
  select(domain_category, web_URL, violation_name, violation_score) |>
  head(2)

## # A tibble: 2 x 4
##   domain_category web_URL           violation_name violation_score
##   <chr>          <chr>             <chr>                  <dbl>
## 1 News and Media https://arstechnica.com/health/ color-contras~      4
## 2 News and Media https://arstechnica.com/health/ color-contrast        4

```

AGGREGATE DATA

Here we aggregated data per page for potentially useful metrics to include in individual summaries we will input into tableau and graph. (We didn't end up using the aggregated values calculated here.)

```

page_summary <- data_clean_deduped |>
  group_by(web_URL, domain) |>
  summarise(
    n_violations = n(),
    avg_severity = mean(violation_score, na.rm = TRUE),
    max_severity = max(violation_score, na.rm = TRUE),
    min_severity = min(violation_score, na.rm = TRUE),
    n_distinctViolation_types = n_distinct(violation_name),
    n_distinctViolation_categories = n_distinct(violation_category),
    .groups = "drop"
  )

```

DOMAIN-LEVEL SUMMARY

```

domain_summary <- page_summary |>
  group_by(domain) |>
  summarise(
    n_pages = n(),
    total_violations = sum(n_violations, na.rm = TRUE),
    avg_violations_per_page = mean(n_violations, na.rm = TRUE),
    avg_severity = mean(avg_severity, na.rm = TRUE),
    .groups = "drop"
  )

```

```

domain_summary |>
  head(4)

```

```

## # A tibble: 4 x 5
##   domain      n_pages total_violations avg_violations_per_page avg_severity
##   <chr>       <int>        <int>                  <dbl>        <dbl>
## 1 3dcart.com      1            4                      4          4.5
## 2 abc.net.au      1            4                      4          3.75
## 3 abcnews.go.com  1           13                     13          4.08
## 4 academia.edu     1            8                      8          3.5

domain_category_summary <- data_clean |>
  group_by(domain_category) |>
  summarise(
    total_violations = n(),
    avg_severity = mean(violation_score, na.rm = TRUE),
    .groups = "drop"
  )

domain_category_summary |>
  head(4)

## # A tibble: 4 x 3
##   domain_category      total_violations avg_severity
##   <chr>                   <int>        <dbl>
## 1 E-commerce              345         3.56
## 2 Educational Platforms    575         3.63
## 3 Government and Public Services 380         3.40
## 4 Health and Wellness     165         3.54

```

VIOLATION-LEVEL SUMMARIES

```

violation_summary <- data_clean |>
  count(violation_name, sort = TRUE) |>
  rename(violations = n)

violation_type_by_domain_category <- data_clean |>
  group_by(domain_category, violation_name) |>
  summarise(violations = n(),
            .groups = "drop")

violation_category_by_domain_category <- data_clean |>
  filter(!is.na(violation_category)) |>
  group_by(domain_category, violation_category) |>
  summarise(violations = n(),
            .groups = "drop")

violation_category_summary <- data_clean |>
  filter(is.na(violation_category)) |>
  count(violation_category, sort = TRUE) |>
  rename(violations = n)

```

```

topViolations <- violation_summary |>
  slice_max(violations, n = 10) |>
  pull(violation_name)

violation_summary_grouped <- violation_summary |>
  mutate(
    violation_group = if_else(
      violation_name %in% topViolations,
      violation_name,
      "Other"
    )
  ) |>
  group_by(violation_group) |>
  summarise(violations = sum(violations), .groups = "drop")

```

```

violation_summary |>
  head(4)

```

```

## # A tibble: 4 x 2
##   violation_name     violations
##   <chr>                  <int>
## 1 color-contrast-enhanced     502
## 2 region                   382
## 3 color-contrast             245
## 4 duplicate-id              230

```

```

violation_type_by_domain_category |>
  head(4)

```

```

## # A tibble: 4 x 3
##   domain_category violation_name     violations
##   <chr>           <chr>                  <int>
## 1 E-commerce       aria-allowed-attr      12
## 2 E-commerce       aria-allowed-role      5
## 3 E-commerce       aria-command-name     2
## 4 E-commerce       aria-dialog-name      2

```

```

violation_category_by_domain_category |>
  head(4)

```

```

## # A tibble: 4 x 3
##   domain_category violation_category violations
##   <chr>           <chr>                  <int>
## 1 E-commerce       Layout                  73
## 2 E-commerce       Syntax                 272
## 3 Educational Platforms Layout                159
## 4 Educational Platforms Semantic            26

```

```

violation_category_summary |>
  head(4)

```

```

## # A tibble: 3 x 2
##   violation_category violations
##   <chr>                <int>
## 1 Syntax                 2634
## 2 Layout                  853
## 3 Semantic                  33

violation_summary_grouped |>
  head(4)

## # A tibble: 4 x 2
##   violation_group     violations
##   <chr>                  <int>
## 1 Other                   1151
## 2 color-contrast            245
## 3 color-contrast-enhanced    502
## 4 duplicate-id              230

```

EXPORT EVERYTHING FOR TABLEAU

```

write.csv(domain_summary, "processed_data/domain_summary.csv", row.names = FALSE)
write.csv(domain_category_summary, "processed_data/domain_category_summary.csv", row.names = FALSE)
write.csv(violation_summary, "processed_data/violation_summary.csv", row.names = FALSE)
write.csv(violation_summary_grouped, "processed_data/violation_summary_grouped.csv", row.names = FALSE)
write.csv(violation_category_summary, "processed_data/violation_category_summary.csv", row.names = FALSE)
write.csv(violation_type_by_domain_category, "processed_data/violation_type_by_domain.csv", row.names =
write.csv(violation_category_by_domain_category, "processed_data/violation_category_by_domain.csv", row

```

finished :D